

Machine Learning Group Project 2023

Ashkan Lotfollahzadeh - Michelle Lee Kruschke - Marco Mancuso Morini - Melissa Nieto
Ridoutt- João Victor Pereira Dos Santos

Introduction

We chose the Kaggle competition <https://www.kaggle.com/competitions/instacart-market-basket-analysis>. Instacart is an American company that offers grocery shopping and delivery services via website and a mobile app. It connects users with a personal shopper in their area to shop and deliver groceries from local participating stores. Although the original objective of the competition was quite complex, we chose to evaluate the problem looking at whether a product would be reordered based on various features of past history, and additionally a Market Basket Analysis to determine which products could possibly be purchased together in the future. The datasets used in the evaluation were provided in 2017, by Instacart, in open sourced data taken from over 200,000 Instacart users.

```
In [1]: #Importing Libraries:

import warnings
warnings.filterwarnings('ignore')

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import f1_score, recall_score, precision_score
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.linear_model import Perceptron
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.svm import SVC
from sklearn.model_selection import RandomizedSearchCV
from mlxtend.frequent_patterns import apriori
from mlxtend.frequent_patterns import association_rules
from xgboost import XGBClassifier
from mlxtend.frequent_patterns import association_rules
from mlxtend.frequent_patterns import apriori
from sklearn.preprocessing import MinMaxScaler
```

```
In [2]: #Creating relevant variables
rs = 42 #Random_state
np.random.seed(rs)
cv = 3
%matplotlib inline
```

```
In [3]: aisles=pd.read_csv('aisles.csv')
departments=pd.read_csv('departments.csv')
products=pd.read_csv('products.csv')
orders=pd.read_csv('orders.csv')
order_products_prior = pd.read_csv('order_products_prior.csv')
order_products_train = pd.read_csv('order_products_train.csv')
```

0. Analysing the data

Aisles

```
In [4]: aisles.head()
```

```
Out[4]:
```

| | aisle_id | aisle |
|---|----------|----------------------------|
| 0 | 1 | prepared soups salads |
| 1 | 2 | specialty cheeses |
| 2 | 3 | energy granola bars |
| 3 | 4 | instant foods |
| 4 | 5 | marinades meat preparation |

```
In [5]: aisles.shape, len(aisles['aisle_id'].unique())
```

```
Out[5]: ((134, 2), 134)
```

```
In [6]: aisles.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 134 entries, 0 to 133
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   aisle_id    134 non-null   int64
1   aisle       134 non-null   object
dtypes: int64(1), object(1)
memory usage: 2.2+ KB
```

Departments

```
In [7]: departments.head()
```

Out[7]:

| | department_id | department |
|---|---------------|------------|
| 0 | 1 | frozen |
| 1 | 2 | other |
| 2 | 3 | bakery |
| 3 | 4 | produce |
| 4 | 5 | alcohol |

In [8]: departments.shape, len(departments['department_id'].unique())

Out[8]: ((21, 2), 21)

In [9]: departments.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21 entries, 0 to 20
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   department_id    21 non-null    int64
1   department       21 non-null    object
dtypes: int64(1), object(1)
memory usage: 464.0+ bytes
```

Products

In [10]: products.head()

Out[10]:

| | product_id | product_name | aisle_id | department_id |
|---|------------|---|----------|---------------|
| 0 | 1 | Chocolate Sandwich Cookies | 61 | 19 |
| 1 | 2 | All-Seasons Salt | 104 | 13 |
| 2 | 3 | Robust Golden Unsweetened Oolong Tea | 94 | 7 |
| 3 | 4 | Smart Ones Classic Favorites Mini Rigatoni Wit... | 38 | 1 |
| 4 | 5 | Green Chile Anytime Sauce | 5 | 13 |

In [11]: products.shape, len(products['product_id'].unique())

Out[11]: ((49688, 4), 49688)

In [12]: products.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 49688 entries, 0 to 49687
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0   product_id      49688 non-null  int64
1   product_name    49688 non-null  object
2   aisle_id        49688 non-null  int64
3   department_id   49688 non-null  int64
dtypes: int64(3), object(1)
memory usage: 1.5+ MB
```

Orders

```
In [13]: orders.head()
```

```
Out[13]:
```

| | order_id | user_id | eval_set | order_number | order_dow | order_hour_of_day | days_since_prior_order |
|---|----------|---------|----------|--------------|-----------|-------------------|------------------------|
| 0 | 2539329 | 1 | prior | 1 | 2 | 8 | NaN |
| 1 | 2398795 | 1 | prior | 2 | 3 | 7 | 15.0 |
| 2 | 473747 | 1 | prior | 3 | 3 | 12 | 21.0 |
| 3 | 2254736 | 1 | prior | 4 | 4 | 7 | 29.0 |
| 4 | 431534 | 1 | prior | 5 | 4 | 15 | 28.0 |

```
In [14]: orders.shape, len(orders['order_id'].unique())
```

```
Out[14]: ((3421083, 7), 3421083)
```

```
In [15]: orders['eval_set'].value_counts()
```

```
Out[15]: prior      3214874  
train      131209  
test       75000  
Name: eval_set, dtype: int64
```

Order_Products_Prior

```
In [16]: order_products_prior.head()
```

```
Out[16]:
```

| | order_id | product_id | add_to_cart_order | reordered |
|---|----------|------------|-------------------|-----------|
| 0 | 2 | 33120 | 1 | 1 |
| 1 | 2 | 28985 | 2 | 1 |
| 2 | 2 | 9327 | 3 | 0 |
| 3 | 2 | 45918 | 4 | 1 |
| 4 | 2 | 30035 | 5 | 0 |

```
In [17]: order_products_prior.shape, len(order_products_prior['order_id'].unique())
```

```
Out[17]: ((32434489, 4), 3214874)
```

```
In [18]: order_products_prior.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32434489 entries, 0 to 32434488
Data columns (total 4 columns):
#   Column                Dtype
---  ---
0   order_id              int64
1   product_id            int64
2   add_to_cart_order     int64
3   reordered              int64
dtypes: int64(4)
memory usage: 989.8 MB
```

Order_Products_Train

```
In [19]: order_products_train.head()
```

```
Out[19]:
```

| | order_id | product_id | add_to_cart_order | reordered |
|---|----------|------------|-------------------|-----------|
| 0 | 1 | 49302 | 1 | 1 |
| 1 | 1 | 11109 | 2 | 1 |
| 2 | 1 | 10246 | 3 | 0 |
| 3 | 1 | 49683 | 4 | 0 |
| 4 | 1 | 43633 | 5 | 1 |

```
In [20]: order_products_train.shape, len(order_products_train['order_id'].unique())
```

```
Out[20]: ((1384617, 4), 131209)
```

```
In [21]: order_products_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1384617 entries, 0 to 1384616
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   order_id              1384617 non-null  int64
1   product_id            1384617 non-null  int64
2   add_to_cart_order     1384617 non-null  int64
3   reordered              1384617 non-null  int64
dtypes: int64(4)
memory usage: 42.3 MB
```

1. Merging the data

```
In [4]: #Creating a unified product dataframe
products=pd.merge(products, aisles, on='aisle_id')
products=pd.merge(products, departments, on='department_id')
```

```
In [5]: #Concating the orders_products train and prior df's:
order_prod_pt = pd.concat([order_products_prior, order_products_train])
```

```
In [6]: #Merging the new df with the prior and train data with the order df:
full_orders_df=pd.merge(orders, order_prod_pt, on='order_id')
```

```
In [7]: #Final dataframe with all datasets merged  
df=pd.merge(full_orders_df, products, on='product_id')
```

```
In [26]: df.head(25)
```

Out[26]:

| | order_id | user_id | eval_set | order_number | order_dow | order_hour_of_day | days_since_prior_ord |
|----|----------|---------|----------|--------------|-----------|-------------------|----------------------|
| 0 | 2539329 | 1 | prior | 1 | 2 | 8 | Na |
| 1 | 2398795 | 1 | prior | 2 | 3 | 7 | 15 |
| 2 | 473747 | 1 | prior | 3 | 3 | 12 | 21 |
| 3 | 2254736 | 1 | prior | 4 | 4 | 7 | 29 |
| 4 | 431534 | 1 | prior | 5 | 4 | 15 | 28 |
| 5 | 3367565 | 1 | prior | 6 | 2 | 7 | 19 |
| 6 | 550135 | 1 | prior | 7 | 1 | 9 | 20 |
| 7 | 3108588 | 1 | prior | 8 | 1 | 14 | 14 |
| 8 | 2295261 | 1 | prior | 9 | 1 | 16 | 0 |
| 9 | 2550362 | 1 | prior | 10 | 4 | 8 | 30 |
| 10 | 1187899 | 1 | train | 11 | 4 | 8 | 14 |
| 11 | 2968173 | 15 | prior | 15 | 1 | 9 | 7 |
| 12 | 1870022 | 15 | prior | 17 | 2 | 16 | 8 |
| 13 | 1911383 | 15 | prior | 18 | 2 | 11 | 7 |
| 14 | 2715276 | 15 | prior | 21 | 1 | 9 | 7 |
| 15 | 487368 | 15 | prior | 22 | 1 | 10 | 14 |
| 16 | 2293453 | 19 | prior | 2 | 5 | 14 | 6 |
| 17 | 1973799 | 19 | prior | 5 | 6 | 12 | 8 |
| 18 | 532817 | 19 | prior | 7 | 4 | 17 | 6 |
| 19 | 1573906 | 21 | prior | 10 | 3 | 10 | 6 |
| 20 | 1593000 | 31 | prior | 10 | 3 | 8 | 7 |
| 21 | 2231262 | 31 | prior | 17 | 3 | 11 | 8 |
| 22 | 2580647 | 43 | prior | 6 | 4 | 16 | 4 |

| | order_id | user_id | eval_set | order_number | order_dow | order_hour_of_day | days_since_prior_ord |
|----|----------|---------|----------|--------------|-----------|-------------------|----------------------|
| 23 | 2187180 | 43 | prior | 9 | 4 | 12 | 3 |
| 24 | 2497897 | 52 | prior | 1 | 1 | 9 | Na |

In [27]: `df.describe()`

| | order_id | user_id | order_number | order_dow | order_hour_of_day | days_since_pr |
|--------------|--------------|--------------|--------------|--------------|-------------------|---------------|
| count | 3.381911e+07 | 3.381911e+07 | 3.381911e+07 | 3.381911e+07 | 3.381911e+07 | 3.17 |
| mean | 1.710566e+06 | 1.029444e+05 | 1.713998e+01 | 2.737285e+00 | 1.343123e+01 | 1.13 |
| std | 9.874008e+05 | 5.946733e+04 | 1.749829e+01 | 2.093296e+00 | 4.246149e+00 | 8.94 |
| min | 1.000000e+00 | 1.000000e+00 | 1.000000e+00 | 0.000000e+00 | 0.000000e+00 | 0.00 |
| 25% | 8.554130e+05 | 5.143500e+04 | 5.000000e+00 | 1.000000e+00 | 1.000000e+01 | 5.00 |
| 50% | 1.710660e+06 | 1.026260e+05 | 1.100000e+01 | 3.000000e+00 | 1.300000e+01 | 8.00 |
| 75% | 2.565587e+06 | 1.544120e+05 | 2.400000e+01 | 5.000000e+00 | 1.600000e+01 | 1.50 |
| max | 3.421083e+06 | 2.062090e+05 | 1.000000e+02 | 6.000000e+00 | 2.300000e+01 | 3.00 |

2. Data Cleaning

In [28]: *#Seeing the information of the dataset*
`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 33819106 entries, 0 to 33819105
Data columns (total 15 columns):
#   Column                                Dtype
---  -
0   order_id                             int64
1   user_id                              int64
2   eval_set                             object
3   order_number                         int64
4   order_dow                            int64
5   order_hour_of_day                    int64
6   days_since_prior_order               float64
7   product_id                           int64
8   add_to_cart_order                    int64
9   reordered                            int64
10  product_name                          object
11  aisle_id                             int64
12  department_id                        int64
13  aisle                                 object
14  department                            object
dtypes: float64(1), int64(10), object(4)
memory usage: 4.0+ GB
```

In [29]: *#Seeing if we have null values:*
`df.isna().sum()`


```
Out[29]: order_id      0
         user_id      0
         eval_set     0
         order_number  0
         order_dow     0
         order_hour_of_day  0
         days_since_prior_order  2078068
         product_id    0
         add_to_cart_order  0
         reordered      0
         product_name    0
         aisle_id        0
         department_id    0
         aisle           0
         department      0
         dtype: int64
```

```
In [8]: #filling nulls with the days_since_prior_order per product mean
df['days_since_prior_order'] = df['days_since_prior_order'].fillna(df.groupby('product_id')['days_since_prior_order'].mean())
```

```
In [31]: #verify the new mean and compare above. The row indexed as 0 had null values in the
y = df.groupby('product_id', as_index=False)['days_since_prior_order'].mean()
print(y['days_since_prior_order'].loc[y['product_id']==196]==df['days_since_prior_order'])
```

```
195      True
Name: days_since_prior_order, dtype: bool
```


```
In [32]: #Seeing if we have null values:
df.isna().sum()
```

```
Out[32]: order_id      0
         user_id      0
         eval_set     0
         order_number  0
         order_dow     0
         order_hour_of_day  0
         days_since_prior_order  40
         product_id    0
         add_to_cart_order  0
         reordered      0
         product_name    0
         aisle_id        0
         department_id    0
         aisle           0
         department      0
         dtype: int64
```

```
In [9]: #we are accepting 40 value=0, it is a particular product but it represents only 0.0
df['days_since_prior_order'] = df['days_since_prior_order'].fillna(0)
df.isna().sum()
```

```
Out[9]: order_id      0
user_id      0
eval_set     0
order_number 0
order_dow    0
order_hour_of_day 0
days_since_prior_order 0
product_id   0
add_to_cart_order 0
reordered    0
product_name 0
aisle_id     0
department_id 0
aisle        0
department   0
dtype: int64
```

```
In [34]: # Graph shows little or no correlation between features
sns.heatmap(df.corr(), annot=False, cmap='Spectral', vmin=-1, vmax=1)
plt.show()
```



3. Reduced Dataset

3.1. Getting a random sample of the dataframe

We need to sample the DataFrame before dividing it into train and validation to be sure that part of the train data is not randomly selected into the validation data.

```
In [11]: # Creating the prior DF so we can divide it:
df_prior = df[df['eval_set'] == 'prior']
```

```
In [36]: # ----- Sampling the data by user' -----
'''
Here it will sample X user_ids and the number of rows would be the sum of the rows
users selected
'''
user_ids_to_sample = list(df_prior['user_id'].sample(300, random_state=rs))

# filter the dataframe to only contain rows with the selected order_ids
prior_red = df_prior[df_prior['user_id'].isin(user_ids_to_sample)]
```

```
In [37]: prior_red.shape, prior_red['user_id'].nunique()
```

```
Out[37]: ((134610, 15), 300)
```

```
In [38]: #dropping columns that we fill that are not important as the others:
prior_red.drop(columns=['order_id', 'user_id', 'aisle_id', 'department_id', 'product_id'], inplace=True)
```

3.2. Dividing the dataset into 3 (divide the dataset into 3 using the eval_set column:

- Prior --> Train
- Train --> Validation
- Test --> Test

```
In [39]: #Dividing the new prior df reduced into train and validation:
train_data, val_data = train_test_split(prior_red, test_size=0.2)
```

```
In [40]: #Test dataset:
df_test = df[df['eval_set'] == 'train'].iloc[0:13696]
```

Dropping the *eval_set* column

Makes no sense to maintain the eval column after the division, so we dropt it.

```
In [41]: train_data.drop(['eval_set'], axis=1, inplace=True);
```

```
In [42]: val_data.drop(['eval_set'], axis=1, inplace=True)
```

```
In [43]: df_test.drop(columns=['eval_set', 'order_id', 'user_id', 'product_id', 'aisle', 'product_name'], inplace=True)
```

```
In [44]: train_data.shape, val_data.shape
```

```
Out[44]: ((107688, 6), (26922, 6))
```

```
In [45]: df_test.shape
```

```
Out[45]: (13696, 6)
```

Dealing with NaNs

```
In [46]: train_data.isna().sum()
```

```
Out[46]: order_number      0
order_dow      0
order_hour_of_day      0
days_since_prior_order      0
add_to_cart_order      0
reordered      0
dtype: int64
```

```
In [47]: val_data.isna().sum()
```

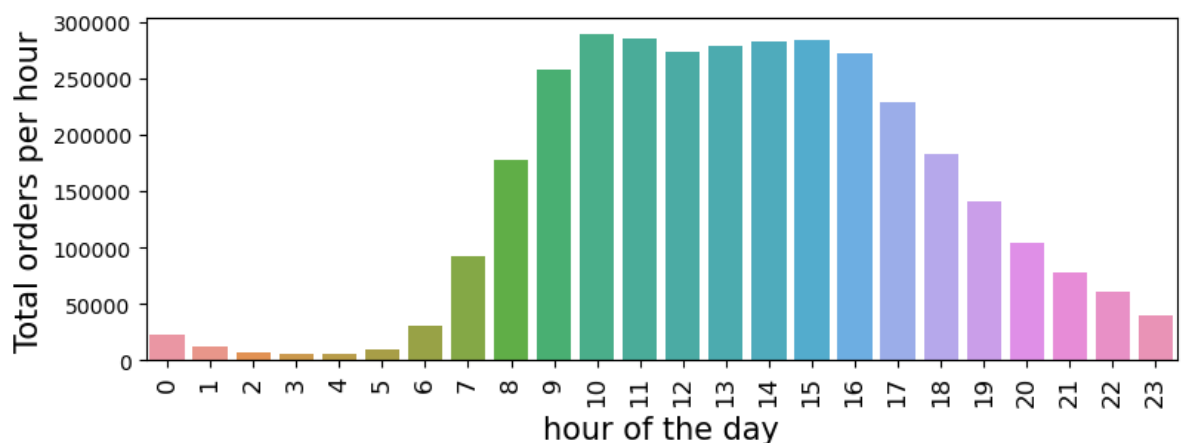
```
Out[47]: order_number      0
order_dow      0
order_hour_of_day      0
days_since_prior_order      0
add_to_cart_order      0
reordered      0
dtype: int64
```

4. Baseline Model


4.1. Data Visualization

Does our sampled data replicate the data distribution of the complete data?

```
In [48]: #order_hour_of_day distribution from the original orders:
orders_hour = orders.groupby('order_hour_of_day')['order_number'].count().to_frame()
plt.figure(figsize=(9,3))
sns.barplot(orders_hour.index, orders_hour.number_of_orders_each_hour, order=orders_hour.index)
plt.xlabel('hour of the day', size=15)
plt.ylabel('Total orders per hour', size=15)
plt.xticks(size=12, rotation=90)
plt.show();
```



```
In [49]: #order_hour_of_day distribution from the train data (sampled):
train_hour = train_data.groupby('order_hour_of_day')['order_number'].count().to_frame()
plt.figure(figsize=(9,3))
sns.barplot(train_hour.index, train_hour.number_of_orders_each_hour, order=train_hour.index)
plt.xlabel('hour of the day', size=15)
plt.ylabel('Total orders per hour', size=15)
plt.xticks(size=12, rotation=90)
plt.show();
```



As both graphs look very similar we concluded that our random sampling replicates pretty well the overall distribution and proportion of the original data.

4.2. Preprocessing

```
In [50]: train_data['days_since_prior_order'] = train_data['days_since_prior_order'].astype
```

```
In [51]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 107688 entries, 25185383 to 30213723
Data columns (total 6 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   order_number                          107688 non-null  int64
1   order_dow                             107688 non-null  int64
2   order_hour_of_day                     107688 non-null  int64
3   days_since_prior_order                 107688 non-null  int8
4   add_to_cart_order                     107688 non-null  int64
5   reordered                             107688 non-null  int64
dtypes: int64(5), int8(1)
memory usage: 5.0 MB
```

```
In [52]: train_data.columns
```

```
Out[52]: Index(['order_number', 'order_dow', 'order_hour_of_day',
               'days_since_prior_order', 'add_to_cart_order', 'reordered'],
              dtype='object')
```

4.3. Splitting the data into the predicting values X and the class y

```
In [53]: #Splitting the train data:
X_train = train_data.drop('reordered', axis=1)
y_train = train_data['reordered']

#Splitting the validation data:
X_val = val_data.drop('reordered', axis=1)
y_val = val_data['reordered']

#Splitting the Test data:
X_test = df_test.drop('reordered', axis=1)
y_test = df_test['reordered']
```

```
In [54]: #Preprocessing this Data:
scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train, y_train)
X_val = scaler.transform(X_val)
X_test = scaler.transform(X_test)
```

4.4. Running Models

4.4.1 Decision Tree Model

We know that a decision tree is too simple of a model, but we wanted to begin with the simplest one before trying random forest and XGBoost.

```
In [55]: #Creating and fitting the Decision Tree model:
model_tree = DecisionTreeClassifier(criterion="entropy", random_state = rs)
model_tree.fit(X_train, y_train);
```

```
In [56]: #Predicting:
y_tree = model_tree.predict(X_val)

#Seeing the accuracy:
accuracy_ho = accuracy_score(y_val, y_tree) * 100

#Understanding the max_depth and preparing a variable for a loop validation:
max_depth = model_tree.tree_.max_depth
parameter_values = range(1,max_depth+1)

print("The accuracy on test set is {0:.1f}%".format(accuracy_ho))
print("The maximum depth of the tree fitted on X_train is {}".format(max_depth))
```

The accuracy on test set is 72.9%

The maximum depth of the tree fitted on X_train is 38

----- Observation -----

We were surprised by the accuracy of the Decision Tree Classifier. Maybe this is not a complex problem as we thought. Maybe the other more complex models will tend to overfit. Still, we will try them anyway. But first let's try to improve the Decision Tree model a little bit more.


```
In [57]: avg_scores = []
for par in parameter_values:
    estimator = DecisionTreeClassifier(criterion="entropy"
                                       , max_depth = par
                                       , random_state = rs
                                       )
    scores = cross_val_score(estimator, X_train, y_train
                             , scoring='accuracy', cv = cv)
    # cross_val_score produces an array with one score for each fold
    avg_scores.append(np.mean(scores))
print(avg_scores)
```

```
[0.7499628556570834, 0.7499628556570834, 0.7520615110318699, 0.7559616670381102,
0.7568624173538371, 0.7575774459549812, 0.7574845850976896, 0.7568345590966495, 0.
7580138919842508, 0.7575495876977936, 0.7576053042121685, 0.757577445954981, 0.756
3888269816507, 0.7540115890349899, 0.7503342990862492, 0.7475577594532353, 0.74438
19181338682, 0.7409739246712724, 0.7364516009211797, 0.732393581457544, 0.72955203
9224426, 0.7276391055642226, 0.7252897258747493, 0.7219931654409034, 0.71991308223
75753, 0.7196159274942427, 0.7186780328355992, 0.7179630042344551, 0.7191701953792
439, 0.7190494762647649, 0.7186223163212242, 0.7189844736646608, 0.718585171978307
7, 0.7187987519500779, 0.7188173241215363, 0.7189751875789318, 0.7188916128073695,
0.7188916128073695]
```

```
In [58]: top_par_cv = parameter_values[np.argmax(avg_scores)]
estimator_best = DecisionTreeClassifier(criterion="entropy", max_depth = top_par_cv)
estimator_best.fit(X_train,y_train);
y_predicted = estimator_best.predict(X_test)
accuracy_cv = accuracy_score(y_test, y_predicted) * 100
print("The accuracy on test set tuned with cross_validation is {:.1f}% with depth 9")
```


The accuracy on test set tuned with cross_validation is 82.1% with depth 9

```
In [59]: #Classification Report on the DT model:
print(classification_report(y_test, y_predicted))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.34 | 0.19 | 0.24 | 2075 |
| 1 | 0.87 | 0.93 | 0.90 | 11621 |
| accuracy | | | 0.82 | 13696 |
| macro avg | 0.60 | 0.56 | 0.57 | 13696 |
| weighted avg | 0.79 | 0.82 | 0.80 | 13696 |

```
In [60]: #Confusion matrix in the DT model:
tree_matrix = confusion_matrix(y_test, y_predicted)
```

```
In [61]: #Plotting the Confusion Matrix in a more vivid and clear way:
plt.rcParams['figure.figsize'] = [4, 4]
disp_tree = ConfusionMatrixDisplay(confusion_matrix=tree_matrix,
                                   display_labels=estimator_best.classes_);
disp_tree.plot();
```



4.5. Fine Tunning

4.5.1. RandomSearchCV

We saw that our simple model was overfitting so we decided to limit the max_depth paramter between 10 to 21. Also, we decided for the RandomSearchCV instead of the GridSearch because the RandomSearchCV could impose less computational power.

```
In [62]: #Definining our models and the params that we will try for each of them:
model_lbls = ['dt' # decision tree
              , 'xgb' # XGBoost
              , 'rf' # Random Forest
              , 'knn' #KNN classifier
              ]

models = {
    'dt': {'name': 'Decision Tree',
           'estimator': DecisionTreeClassifier(random_state=rs),
           'param': [{'criterion': ['gini', 'entropy'],
                        'min_samples_split': [*range(1,10)],
                        'max_depth': [*range(1,max_depth+1)],
                        'class_weight': [None, 'balanced'],
                        'min_samples_leaf': [*range(1,10)]
                       ]}],

    'xgb': {'name': 'XGBoost',
            'estimator': XGBClassifier(random_state=rs),
            'param': [{'n_estimators': [10, 20, 30, 40],
                        'learning_rate': [0.001, 0.005, 0.01],
                        'n_jobs': [None, -1]
                       ]}],

    'rf': {'name': 'Random Forest',
           'estimator': RandomForestClassifier(random_state=rs),
           'param': [{'criterion': ['gini', 'entropy'],
                        'max_depth': [*range(4,max_depth+1)],
                        'n_estimators': [*range(10,60,10)],
                        'min_samples_leaf': [*range(1,10)],
                        'min_samples_split': [*range(1,10)]
                       ]}],

    'knn': {'name': 'K Nearest Neighbor',
            'estimator': KNeighborsClassifier(),
            'param': [{'n_neighbors': [2],
                        'weights': ['uniform', 'distance'],
                        'metric': ['euclidean', 'manhattan', 'cityblock'],
                       ]}]
}
```

```
In [63]: scorings = ['accuracy', 'precision_macro', 'recall_macro', 'f1_macro']
```

```
In [64]: #Running a GridSearch for each model and each different combination of params:
clfs = []
my_models = pd.DataFrame(columns=['scoring', 'model', 'best_params', 'accuracy', 'precision_macro', 'recall_macro', 'f1_macro'])

for model in model_lbls:
    for scoring in scorings:
        clf = RandomizedSearchCV(models[model]['estimator'], models[model]['param'],
                                  scoring=scoring, n_jobs=-1, n_iter=200, verbose=2)
        #Appending the clfs list with each clf fit
```



```

clf.fit(X_train, y_train)
clfs.append(clf)

#creating the accuracy variable
y_pred = clf.predict(X_val)
accuracy = accuracy_score(y_val,y_pred)

cr = classification_report(y_val,y_pred, output_dict=True)

#creating the precision_macro variable
precision_macro = precision_score(y_val, y_pred, average='macro')

#creating the recall_macro variable
recall_macro = recall_score(y_val, y_pred, average='macro')

#creating the f1_macro variable
f1_macro = f1_score(y_val, y_pred, average='macro')

best_params = clf.best_params_

#Appending the dataset row by row
row = [scoring, model, best_params, accuracy, precision_macro, recall_macro, f1_macro]
my_models.loc[len(my_models)] = row

```

```

Fitting 3 folds for each of 200 candidates, totalling 600 fits
Fitting 3 folds for each of 200 candidates, totalling 600 fits
Fitting 3 folds for each of 200 candidates, totalling 600 fits
Fitting 3 folds for each of 200 candidates, totalling 600 fits
Fitting 3 folds for each of 24 candidates, totalling 72 fits
Fitting 3 folds for each of 24 candidates, totalling 72 fits
Fitting 3 folds for each of 24 candidates, totalling 72 fits
Fitting 3 folds for each of 24 candidates, totalling 72 fits
Fitting 3 folds for each of 200 candidates, totalling 600 fits
Fitting 3 folds for each of 200 candidates, totalling 600 fits
Fitting 3 folds for each of 200 candidates, totalling 600 fits
Fitting 3 folds for each of 200 candidates, totalling 600 fits
Fitting 3 folds for each of 6 candidates, totalling 18 fits
Fitting 3 folds for each of 6 candidates, totalling 18 fits
Fitting 3 folds for each of 6 candidates, totalling 18 fits
Fitting 3 folds for each of 6 candidates, totalling 18 fits

```

In [65]: *#Plotting our models dataframe:*
my_models

Out[65]:

| | scoring | model | best_params | accuracy | precision_macro | recall_macro | f1_macro |
|----|-----------------|-------|---|----------|-----------------|--------------|----------|
| 0 | accuracy | dt | {'min_samples_split': 2, 'min_samples_leaf': 2... | 0.758042 | 0.719503 | 0.645197 | 0.658558 |
| 1 | precision_macro | dt | {'min_samples_split': 6, 'min_samples_leaf': 2... | 0.756222 | 0.726250 | 0.630649 | 0.641953 |
| 2 | recall_macro | dt | {'min_samples_split': 3, 'min_samples_leaf': 2... | 0.704257 | 0.665862 | 0.688617 | 0.670509 |
| 3 | f1_macro | dt | {'min_samples_split': 3, 'min_samples_leaf': 3... | 0.751430 | 0.698781 | 0.677464 | 0.685452 |
| 4 | accuracy | xgb | {'n_jobs': None, 'n_estimators': 40, 'learning... | 0.757410 | 0.734108 | 0.627439 | 0.637984 |
| 5 | precision_macro | xgb | {'n_jobs': None, 'n_estimators': 30, 'learning... | 0.758116 | 0.736356 | 0.627757 | 0.638369 |
| 6 | recall_macro | xgb | {'n_jobs': None, 'n_estimators': 40, 'learning... | 0.757410 | 0.734108 | 0.627439 | 0.637984 |
| 7 | f1_macro | xgb | {'n_jobs': None, 'n_estimators': 40, 'learning... | 0.757410 | 0.734108 | 0.627439 | 0.637984 |
| 8 | accuracy | rf | {'n_estimators': 40, 'min_samples_split': 8, '... | 0.783857 | 0.758353 | 0.684369 | 0.702399 |
| 9 | precision_macro | rf | {'n_estimators': 30, 'min_samples_split': 5, '... | 0.752842 | 0.764583 | 0.600422 | 0.599991 |
| 10 | recall_macro | rf | {'n_estimators': 30, 'min_samples_split': 2, '... | 0.770299 | 0.725443 | 0.692570 | 0.703965 |
| 11 | f1_macro | rf | {'n_estimators': 50, 'min_samples_split': 2, '... | 0.775871 | 0.735938 | 0.691082 | 0.705052 |
| 12 | accuracy | knn | {'weights': 'distance', 'n_neighbors': 2, 'met... | 0.761422 | 0.714672 | 0.720672 | 0.717446 |
| 13 | precision_macro | knn | {'weights': 'distance', 'n_neighbors': 2, 'met... | 0.761422 | 0.714672 | 0.720672 | 0.717446 |
| 14 | recall_macro | knn | {'weights': 'uniform', 'n_neighbors': 2, 'metr... | 0.737241 | 0.709776 | 0.745475 | 0.714017 |

| | scoring | model | best_params | accuracy | precision_macro | recall_macro | f1_macro |
|----|----------|-------|---|----------|-----------------|--------------|----------|
| 15 | f1_macro | knn | {'weights': 'distance', 'n_neighbors': 2, 'metric': 'manhattan'} | 0.761422 | 0.714672 | 0.720672 | 0.717446 |

In [66]: *#Seeing our best model for each score type:*

```
print("Best Accuracy model is {} using this best parameters: {}".format(my_models['accuracy'], my_models['best_params']))
print('-'*90)

print("Best Precision model is {} using this best parameters: {}".format(my_models['precision_macro'], my_models['best_params']))
print('-'*90)

print("Best Recall model is {} using this best parameters: {}".format(my_models['recall_macro'], my_models['best_params']))
print('-'*90)

print("Best F1 model is {} using this best parameters: {}".format(my_models['f1_macro'], my_models['best_params']))
```

Best Accuracy model is rf using this best parameters: {'n_estimators': 40, 'min_samples_split': 8, 'min_samples_leaf': 1, 'max_depth': 32, 'criterion': 'entropy'}.

Best Precision model is rf using this best parameters: {'n_estimators': 30, 'min_samples_split': 5, 'min_samples_leaf': 6, 'max_depth': 4, 'criterion': 'entropy'}.

Best Recall model is knn using this best parameters: {'weights': 'uniform', 'n_neighbors': 2, 'metric': 'manhattan'}.

Best F1 model is knn using this best parameters: {'weights': 'distance', 'n_neighbors': 2, 'metric': 'manhattan'}.

4.6. Testing the best model

Testing the best model to see if it can reproduce the same accuracy, precision, etc.

4.6.1. Random Forest (The Theoretically Best model)

```
In [67]: best_model = RandomForestClassifier(criterion="entropy",
                                           random_state = rs,
                                           max_depth = 32,
                                           n_estimators=40,
                                           min_samples_split=8,
                                           min_samples_leaf=1)

best_model.fit(X_train, y_train)
```

Out[67]: RandomForestClassifier(criterion='entropy', max_depth=32, min_samples_split=8, n_estimators=40, random_state=42)

```
In [68]: #Predicting:
y_best = best_model.predict(X_test)

#Seeing the accuracy:
accuracy_best = accuracy_score(y_test, y_best) * 100
```

```
#creating the precision_macro variable
precision_best = precision_score(y_test, y_best, average='macro') * 100

#creating the recall_macro variable
recall_best = recall_score(y_test, y_best, average='macro') * 100

#creating the f1_macro variable
f1_best = f1_score(y_test, y_best, average='macro') * 100

print("The accuracy on test set is {:.1f}%".format(accuracy_best))
```

The accuracy on test set is 80.4%


```
In [69]: #Classification Report on the RF model:
print(classification_report(y_test, y_best))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.31 | 0.24 | 0.27 | 2075 |
| 1 | 0.87 | 0.90 | 0.89 | 11621 |
| accuracy | | | 0.80 | 13696 |
| macro avg | 0.59 | 0.57 | 0.58 | 13696 |
| weighted avg | 0.79 | 0.80 | 0.79 | 13696 |

```
In [70]: print(confusion_matrix(y_test, y_best))
```

```
[[ 502 1573]
 [1109 10512]]
```

```
In [71]: #Plotting the Confusion Matrix in a more vivid and clear way:
plt.rcParams['figure.figsize'] = [4, 4]
disp = ConfusionMatrixDisplay(confusion_matrix=confusion_matrix(y_test, y_best),
                              display_labels=best_model.classes_)
disp.plot();
```



Market Basket Analysis Association and Column creation

We need to sample the dataframe again because otherwise our notebook wouldn't run the code below.

```
In [87]: # ----- Sampling the data by user' -----
# create a list of order_ids to sample the data
'''
Here it will sample X user_ids and the number of rows would be the sum of the rows
users selected
'''
user_ids_to_sample2 = list(df_prior['user_id'].sample(15000, random_state=rs))

# filter the dataframe to only contain rows with the selected order_ids
filtered_df2 = df_prior[df_prior['user_id'].isin(user_ids_to_sample2)]

# group the filtered dataframe by order_id and apply the sample function to each group
sampled_groups2 = [group.sample(n=len(group), replace=True) for _, group in filtered_df2.groupby('order_id')]

# concatenate the sampled groups into a single dataframe
sampled_df2 = pd.concat(sampled_groups2)

# reset the index of the sampled dataframe
prior_red2 = sampled_df2.reset_index(drop=True)
```

```
In [88]: prior_red2.shape
```

```
Out[88]: (5380527, 15)
```

```
In [89]: #Group the products by order_id and user_id to create transactions:
transactions = prior_red2.groupby(['order_id', 'user_id'])['product_id'].apply(list)
```

```
In [90]: #Convert the products into binary indicators for each transaction:
```

```
from mlxtend.preprocessing import TransactionEncoder

te = TransactionEncoder()
te_ary = te.fit_transform(transactions['products'])
transactions = pd.DataFrame(te_ary, columns=te.columns_)
```

```
In [98]: #Use the Apriori algorithm to find frequent itemsets:
```

```
from mlxtend.frequent_patterns import apriori

frequent_itemsets = apriori(transactions, min_support=0.01, use_colnames=True)
```

```
In [99]: frequent_itemsets
```

Out[99]:

| | support | itemsets |
|-----|----------|----------------|
| 0 | 0.016407 | (4605) |
| 1 | 0.010004 | (4799) |
| 2 | 0.021629 | (4920) |
| 3 | 0.015997 | (5077) |
| 4 | 0.011712 | (5450) |
| ... | ... | ... |
| 79 | 0.019091 | (49235) |
| 80 | 0.023270 | (49683) |
| 81 | 0.012566 | (13176, 21137) |
| 82 | 0.013100 | (13176, 47209) |
| 83 | 0.011019 | (21137, 24852) |

84 rows × 2 columns

```
In [100... #Generate association rules:
from mlxtend.frequent_patterns import association_rules

rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1)
```

In [101... rules

Out[101]:

| | antecedents | consequents | antecedent support | consequent support | support | confidence | lift | leverage |
|---|-------------|-------------|--------------------|--------------------|----------|------------|----------|----------|
| 0 | (13176) | (21137) | 0.096375 | 0.075941 | 0.012566 | 0.130384 | 1.716914 | 0.005247 |
| 1 | (21137) | (13176) | 0.075941 | 0.096375 | 0.012566 | 0.165468 | 1.716914 | 0.005247 |
| 2 | (13176) | (47209) | 0.096375 | 0.060643 | 0.013100 | 0.135930 | 2.241465 | 0.007256 |
| 3 | (47209) | (13176) | 0.060643 | 0.096375 | 0.013100 | 0.216021 | 2.241465 | 0.007256 |
| 4 | (21137) | (24852) | 0.075941 | 0.116107 | 0.011019 | 0.145100 | 1.249705 | 0.002202 |
| 5 | (24852) | (21137) | 0.116107 | 0.075941 | 0.011019 | 0.094903 | 1.249705 | 0.002202 |

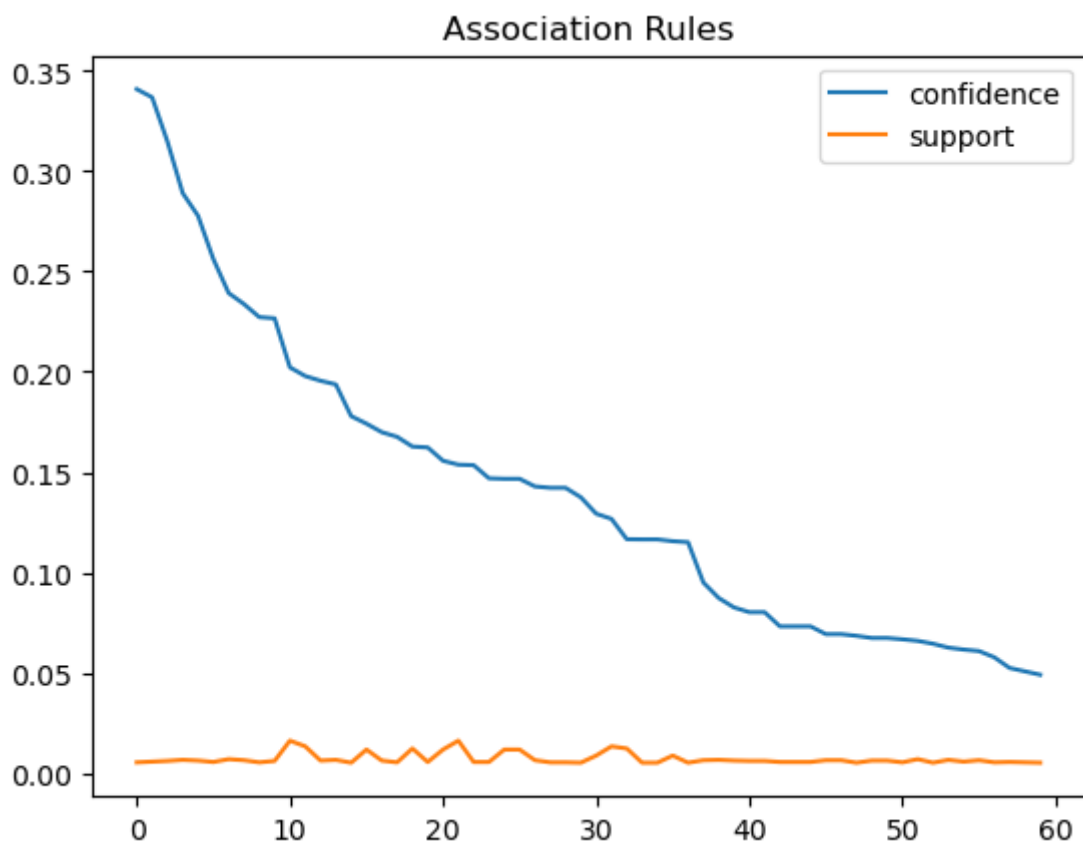
- **antecedent support:** the proportion of transactions that contain the antecedent set of items.
- **consequent support:** the proportion of transactions that contain the consequent set of items.
- **support:** the proportion of transactions that contain both the antecedent and consequent sets of items.
- **confidence:** the proportion of transactions that contain both the antecedent and consequent sets of items among the transactions that contain the antecedent set of

items. It measures how often the consequent items appear in transactions that contain the antecedent items.

- **lift**: the ratio of observed support to expected support if the antecedent and consequent were independent. A lift value greater than 1 indicates a positive correlation between the antecedent and consequent, while a value less than 1 indicates a negative correlation.
- **leverage**: the difference between observed support and expected support if the antecedent and consequent were independent. It measures the deviation of the observed frequency of both sets of items from what would be expected if there was no association.
- **conviction**: the ratio of the expected frequency of the antecedent to the frequency of the antecedent that are not followed by the consequent. It measures the degree of implication between the antecedent and consequent.
- **zhangs_metric**: a metric that takes into account the direction and strength of association between the antecedent and consequent. A value greater than 0 indicates a positive correlation, while a value less than 0 indicates a negative correlation.

In [63]: *# Graph representing confidence and support*

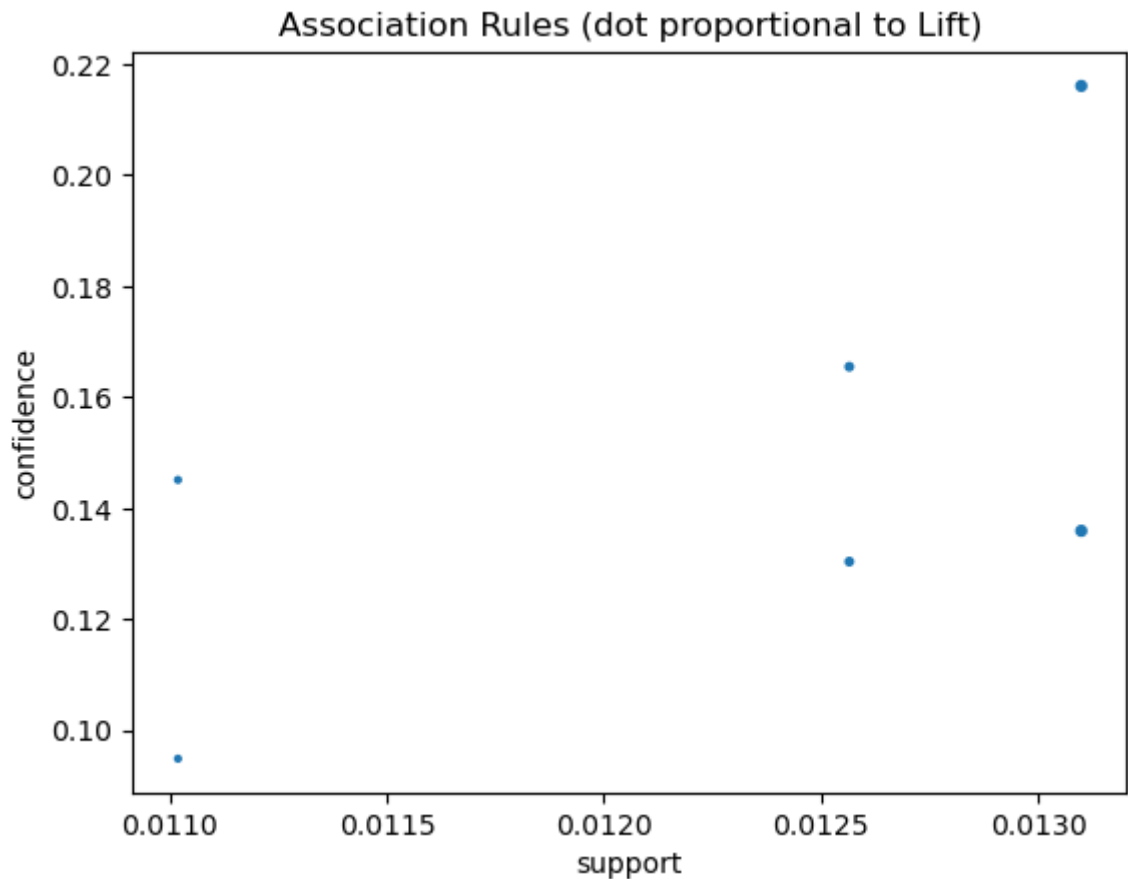
```
sorted_rules=rules.sort_values(by=['confidence','support'],ascending=False).reset_  
sorted_rules[['confidence','support']].plot(title='Association Rules');
```



In [105...]

```
# 1.8 is chosen empirically to obtain the best graphical effect  
s = [3*n for n in rules.lift]  
rules.plot.scatter(x='support',  
                  y='confidence',
```

```
title='Association Rules (dot proportional to Lift)',  
s=s);
```



Conclusions

We are able to predict the reordering with a good performance on the sampled data. Unfortunately due to computational problems we were not able to make predictions on the whole dataset: it would be nice to be able to evaluate the performance on the whole dataset. In fact we have an F1 score of 89% and an accuracy of 87% on the test set. Which suggests that we are likely to be able to predict with good performance on the entire dataset as well. As far as the association rule is concerned, we have obtained interesting rules, but here too, the impossibility of being able to evaluate the performance on the entire dataset leads us to limited conclusions.

In []:

In []:

In []:

In []:

In []: